## Gaussian Naive Bayes (GNB)

# Import required libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model\_selection import train\_test\_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import accuracy\_score, classification\_report, confusion\_ma

url = 'https://github.com/manas8173/diabeties/raw/main/heart\_data.csv'
heart = pd.read\_csv(url, sep=',') # explicitly use comma separator
display(heart)

id         age         gender         height         weight         ap_hi         ap_hi         ap_lo         cholesterol         gluc         smoke         ap_hi           0         0         18393         2         168         62.0         110         80         1         1         0           1         1         20228         1         156         85.0         140         90         3         1         0           2         2         18857         1         165         64.0         130         70         3         1         0           3         3         17623         2         169         82.0         150         100         1         1         0           4         4         17474         1         156         56.0         100         60         1         1         0 <td< th=""><th><math>\rightarrow</math></th><th></th><th></th><th></th><th></th><th></th><th></th><th></th><th>_</th><th></th><th>_</th><th></th><th></th></td<>	$\rightarrow$								_		_		
1       1       20228       1       156       85.0       140       90       3       1       0         2       2       18857       1       165       64.0       130       70       3       1       0         3       3       17623       2       169       82.0       150       100       1       1       0         4       4       17474       1       156       56.0       100       60       1       1       0	_		id	age	gender	height	weight	ap_h1	ap_1o	cholesterol	gluc	smoke	ć
2       2       18857       1       165       64.0       130       70       3       1       0         3       3       17623       2       169       82.0       150       100       1       1       0         4       4       17474       1       156       56.0       100       60       1       1       0		0	0	18393	2	168	62.0	110	80	1	1	0	
3       3       17623       2       169       82.0       150       100       1       1       0         4       4       17474       1       156       56.0       100       60       1       1       0  .		1	1	20228	1	156	85.0	140	90	3	1	0	
4       4       17474       1       156       56.0       100       60       1       1       0		2	2	18857	1	165	64.0	130	70	3	1	0	
.		3	3	17623	2	169	82.0	150	100	1	1	0	
69995       99993       19240       2       168       76.0       120       80       1       1       1         69996       99995       22601       1       158       126.0       140       90       2       2       0         69997       99996       19066       2       183       105.0       180       90       3       1       0         69998       99998       22431       1       163       72.0       135       80       1       2       0		4	4	17474	1	156	56.0	100	60	1	1	0	
69996       99995       22601       1       158       126.0       140       90       2       2       0         69997       99996       19066       2       183       105.0       180       90       3       1       0         69998       99998       22431       1       163       72.0       135       80       1       2       0										•••			
69997       99996       19066       2       183       105.0       180       90       3       1       0         69998       99998       22431       1       163       72.0       135       80       1       2       0		69995	99993	19240	2	168	76.0	120	80	1	1	1	
<b>69998</b> 99998 22431 1 163 72.0 135 80 1 2 0		69996	99995	22601	1	158	126.0	140	90	2	2	0	
		69997	99996	19066	2	183	105.0	180	90	3	1	0	
<b>69999</b> 99999 20540 1 170 72.0 120 80 2 1 0		69998	99998	22431	1	163	72.0	135	80	1	2	0	
		69999	99999	20540	1	170	72.0	120	80	2	1	0	

70000 rows × 13 columns

# 🖢 Step 1: Load & Preprocess

from sklearn.model\_selection import train\_test\_split from sklearn.preprocessing import StandardScaler

```
# Assume `heart` is already loaded
X = heart.drop("cardio", axis=1)
```

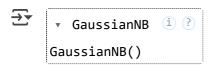
y = heart["cardio"]

```
# Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## Step 2: Train GNB Model

```
from sklearn.naive_bayes import GaussianNB
gnb_model = GaussianNB()
gnb_model.fit(X_train_scaled, y_train)
```



### 📊 Step 3: Evaluate GNB Accuracy

```
from sklearn.metrics import accuracy_score, classification_report

y_pred_gnb = gnb_model.predict(X_test_scaled)

print("Accuracy:", accuracy_score(y_test, y_pred_gnb))
print(classification_report(y_test, y_pred_gnb))
```

### Accuracy: 0.5887857142857142 precision recall f1-score support 0 0.56 0.88 0.68 7004 0.71 0.30 0.42 6996 0.59 14000 accuracy 0.63 0.59 0.55 14000 macro avg 0.59 0.55 14000 weighted avg 0.63

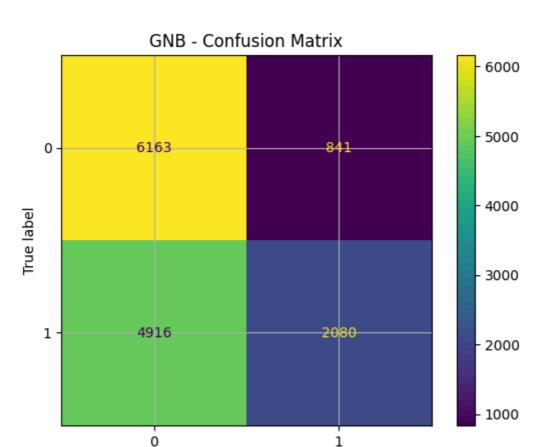
## Step 4: Confusion Matrix

from sklearn.metrics import confusion\_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

```
cm = confusion_matrix(y_test, y_pred_gnb)
```

```
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot()
plt.title("GNB - Confusion Matrix")
plt.grid()
plt.show()
```





Predicted label

### Step 5: ROC Curve

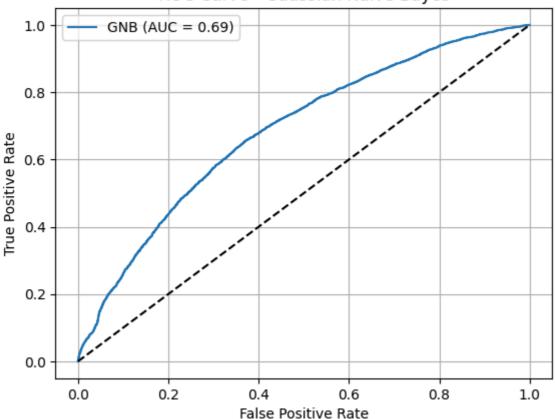
```
from sklearn.metrics import roc_curve, auc

y_prob_gnb = gnb_model.predict_proba(X_test_scaled)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_prob_gnb)
roc_auc = auc(fpr, tpr)

plt.plot(fpr, tpr, label="GNB (AUC = %0.2f)" % roc_auc)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Gaussian Naive Bayes")
plt.legend()
plt.grid()
plt.show()
```

 $\rightarrow$ 

## **ROC Curve - Gaussian Naive Bayes**



Step 6: Train vs Validation Accuracy Curve python Copy Edit

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
train_sizes = np.linspace(0.1, 0.99, 10)
train scores = []
val_scores = []
for frac in train sizes:
   X_part, _, y_part, _ = train_test_split(X_train_scaled, y_train, train_size
    model = GaussianNB()
    model.fit(X_part, y_part)
    train_scores.append(model.score(X_part, y_part))
    val_scores.append(model.score(X_test_scaled, y_test))
plt.figure(figsize=(8, 5))
plt.plot(train_sizes, train_scores, label='Training Accuracy', marker='o')
plt.plot(train_sizes, val_scores, label='Validation Accuracy', marker='s')
plt.xlabel("Simulated Epochs (Training Size Fraction)", fontsize=12)
plt.ylabel("Accuracy", fontsize=12)
plt.title("Simulated Train vs Validation Accuracy - GNB", fontsize=14)
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



# Simulated Train vs Validation Accuracy - GNB Training Accuracy Validation Accuracy 0.62 0.61 0.59

0.6

Simulated Epochs (Training Size Fraction)

0.8

1.0

### Step 2: Train XGBoost with Epoch Tracking

0.2

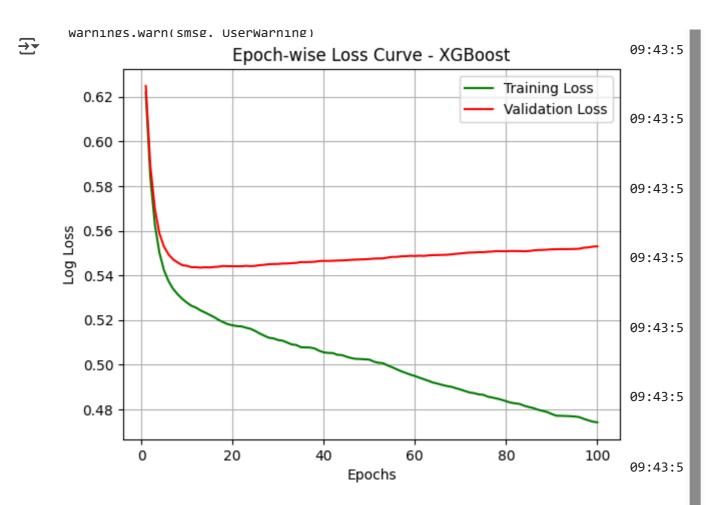
```
from xgboost import XGBClassifier
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
# Define model with eval_metric inside the constructor
xgb model = XGBClassifier(
    n estimators=100,
    use label encoder=False,
    eval_metric='logloss' # ✓ move here
)
# Train and capture evaluation result
xgb_model.fit(
   X_train_scaled, y_train,
    eval_set=[(X_train_scaled, y_train), (X_test_scaled, y_test)],
   verbose=False
)
# Get training results
results = xgb model.evals result()
epochs = range(1, len(results['validation_0']['logloss']) + 1)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:09]
     Parameters: { "use_label_encoder" } are not used.
       warnings.warn(smsg, UserWarning)
```

### Step 3: Plot Epoch-wise Accuracy (Simulated)

```
train_acc = []
val_acc = []
for i in range(1, 101): # from 1 to 100 estimators
    model = XGBClassifier(
        n_estimators=i,
        use_label_encoder=False,
        eval_metric='logloss'
    model.fit(X_train_scaled, y_train)
    train_acc.append(accuracy_score(y_train, model.predict(X_train_scaled)))
    val_acc.append(accuracy_score(y_test, model.predict(X_test_scaled)))
plt.plot(range(1, 101), train_acc, label='Training Accuracy', color='blue')
plt.plot(range(1, 101), val_acc, label='Validation Accuracy', color='orange')
plt.xlabel("Epochs (Boosting Rounds)")
plt.ylabel("Accuracy")
plt.title("Epoch-wise Accuracy - XGBoost")
plt.legend()
plt.grid(True)
plt.show()
```

```
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:4:
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:4
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:4
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:50
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:5
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:5
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:5
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:5
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:5
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:5
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:54]
    Parameters: { "use label encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:54]
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:54]
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:5]
    Parameters: { "use_label_encoder" } are not used.
      warnings.warn(smsg, UserWarning)
    /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:43:5]
    Parameters: { "use_label_encoder" } are not used.
Step 4: Epoch-wise Loss Curve (Direct from XGBoost)
      warnings.warn(smsg, UserWarning)
```

```
plt.plot(epochs, results['validation_0']['logloss'], label='Training Loss', color='green'
plt.plot(epochs, results['validation_1']['logloss'], label='Validation Loss', color='red'
plt.xlabel("Epochs")
plt.ylabel("Log Loss")
plt.title("Epoch-wise Loss Curve - XGBoost")
plt.legend()
plt.grid(True)
plt.show()
```



warnings.warn(smsg, UserWarning)

Steps 5/16 on fusion by that bix3+1R post-packages/xgboost/core.py:158: UserWarning: [09:43:5]

Parameters: { "use label encoder" } are not used.

from sklearn.metrics import confusion\_matrix, ConfusionMatrixDisplay, classification\_repo

```
y_pred_xgb = xgb_model.predict(X_test_scaled)

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred_xgb)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot()
plt.title("Confusion Matrix - XGBoost")
plt.grid()
plt.show()

# Report
print("Classification Report:\n", classification_report(y_test, y_pred_xgb))
```

/usr/local/lib/pvtnon3.11/dist-packages/xgboost/core.pv:158: Userwarning: [09:43:5]  $\overline{\rightarrow}$ Confusion Matrix - XGBoost 5000 rning: [09:43:5 4500 0 5361 rning: [09:43:5 - 4000 Frue label rning: [09:43:5 3500 rning: [09:43:5 4874 1 - 2500 rning: [09:44:0 - 2000 rning: [09:44:0 1 0 Predicted label ClassifigstwonnRemogt:UserWarning) /usr/local/lib/python3om1/dirteplakagas%gboost%upperhy:158: UserWarning: [09:44:00 Parameters: { "use\_label\_encoder" } are not used. 0.74 7004 0.72 0.77 warnings.Warn(sms@,70serWar@in@) 0.72 6996 /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:44:0 Parameturscy{ "use\_label\_encoder" } are0n78 used14000 0.73 0.73 0.73 14000 macro avg 14000 wewghtedgsvgarn(smsg,78serWar@ing) 0.73 /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:44:0 Parameters: { "use label encoder" } are not used. 📊 Stepuán:nRops նահրտա (smsg, UserWarning) /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:44:0 from sklearn.metrics import roc curve, auc y\_prob\_xgb = xgb\_model.predict\_proba(X\_test\_scaled)[:, 1] fpr, tpr, \_ = roc\_curve(y\_test, y\_prob\_xgb) roc auc = auc(fpr, tpr) plt.plot(fpr, tpr, label="XGBoost (AUC = %0.2f)" % roc auc) plt.plot([0, 1], [0, 1], 'k--') plt.xlabel("False Positive Rate") plt.ylabel("True Positive Rate") plt.title("ROC Curve - XGBoost") plt.legend()

```
warnings.warn(smsg, UserWarning)
```

plt.grid() plt.show()

/usr/local/lih/nython3 11/dist-nackages/xghoost/core nv·158. UserWarning. [09.44.0]



